



Shared-Memory Programming: Pthread

National Tsing-Hua University
2019, Summer Semester




Outline

- Shared-memory Programming
- Pthread
- Synchronization Problem & Tools

Shared-Memory Programming


- **Definition:** Processes communicate or work together with each other **through a shared memory space** which can be accessed by all processes
 - **Faster & more efficient than message passing**
- **Many issues as well:**
 - **Synchronization**
 - **Deadlock**
 - **Cache coherence**
- **Programming techniques:**
 - **Parallelizing compiler**
 - **Unix processes**
 - **Threads (**Pthread**, Java)**



Threads vs. Processes

- **Process (heavyweight process):** complete separate program with its own variables, stack, heap, and everything else.
- **Thread (lightweight process):** share the **same memory space** for global variables, resources
- **In Linux:**
 - Threads are created via **clone a process** with a flag to indicate the **level of sharing**

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.



Why Thread?

■ Lower creation/management cost vs. Process

platform	fork()	pthread_create()	speedup
AMD 2.4 GHz Opteron	17.6	1.4	15.6x
IBM 1.5 GHz POWER4	104.5	2.1	49.8x
INTEL 2.4 GHz Xeon	54.9	1.6	34.3x
INTEL 1.4 GHz Itanium2	54.5	2.0	27.3x

■ Faster inter-process communication vs. MPI

platform	MPI Shared Memory BW (GB/sec)	Pthreads Worst Case Memory-to-CPU BW (GB/sec)	speedup
AMD 2.4 GHz Opteron	1.2	5.3	4.4x
IBM 1.5 GHz POWER4	2.1	4	1.9x
INTEL 2.4 GHz Xeon	0.3	4.3	14.3x
INTEL 1.4 GHz Itanium2	1.8	6.4	3.6x

Outline

- Shared-memory Programming
- Pthread
 - What is Pthread
 - Pthread Creation
 - Pthread Joining & Detaching
- Synchronization Problem & Tools

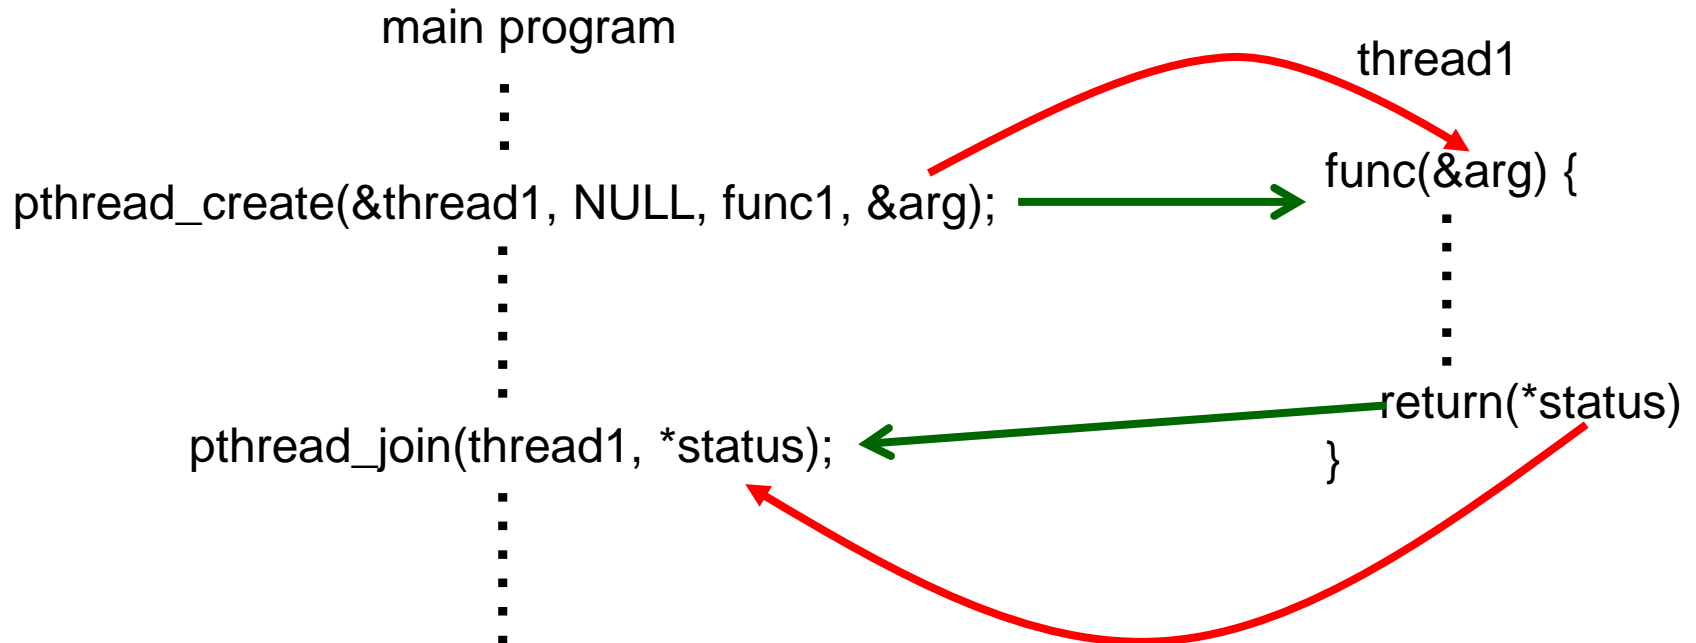
What is Pthread?

- Historically, hardware vendors have implemented their own proprietary versions of threads
- **POSIX** (Potable Operating System Interface) standard is specified for portability across Unix-like systems
 - Similar concept as MPI for message passing libraries
- **Pthread** is the implementation of POSIX standard for thread
 - Same relation between MPICH and MPI

Pthread Creation

■ pthread_create(thread,attr,routine,arg)

- **thread**: An **unique identifier** (token) for the new thread
- **attr**: It is used to set **thread attributes**. NULL for the default values
- **routine**: The routine that the thread will execute once it is created
- **arg**: A **single argument** that may be **passed to routine**



Example

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5

void *PrintHello(void *threadId) {
    int* data = static_cast<int*> (threadId);
    printf("Hello World! It's me, thread #%d!\n", *data);
    pthread_exit(NULL);
}

int main (int argc, char *argv[]) {
    pthread_t threads[NUM_THREADS];
    int tids[NUM_THREADS];
    for(int i=0; i<NUM_THREADS; i++){
        tids[i] = i;
        pthread_create(&threads[i], NULL, PrintHello, (void *)&tids[i]);
    }
    /* Last thing that main() should do */
    pthread_exit(NULL);
}
```

Pthread Joining & Detaching

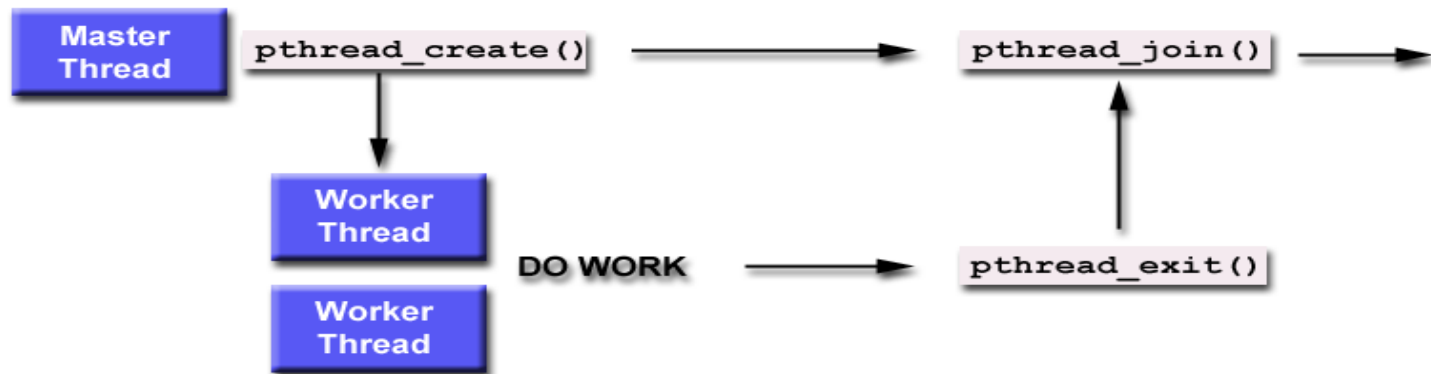
■ pthread_join(threadId, status)

- Blocks until the specified *threadId* thread terminates
- One way to accomplish synchronization between threads
- Example: to create a pthread barrier

```
for (int i=0; i<n; i++) pthread_join(thread[i], NULL);
```

■ pthread_detach(threadId)

- Once a thread is **detached**, it can **never** be joined
- Detach a thread could free some system resources



Outline

- Shared-memory Programming
- Pthread
- Synchronization Problem & Tools
 - Pthread
 - ◆ Mutually exclusion Lock
 - ◆ Condition variable
 - POSIX Semaphore
- Other issues

Synchronization Problem

- The outcome of data content should **NOT** be decided by the **execution order among processes**

- **Instructions** of individual processes/threads may be **interleaved** in time

- E.g.: Assume variable **“counter”** is **shared by processes**

Process0

```
main() {
```

```
...
```

```
    counter++;
```

```
...
```

```
}
```

Process1

```
main() {
```

```
...
```

```
    counter--;
```

```
...
```

```
}
```

- The statement **“counter++”** & **“counter--”** may be implemented in machine language as:

```
move ax, counter
add  ax, 1
move counter, ax
```

```
move bx, counter
sub  bx, 1
move counter, bx
```

Instruction Interleaving

- Assume counter is initially 5. One interleaving of statement is:

producer: move ax, counter ➔ ax = 5

producer: add ax, 1 ➔ ax = 6

context switch

consumer: move bx, counter ➔ bx = 5

consumer: sub bx, 1 ➔ bx = 4

context switch

producer: move counter, ax ➔ counter = 6

context switch

consumer: move counter, bx ➔ counter = 4


- The value of counter may be either 4, 5, or 6
- The ONLY correct result is 5!

Outline

- Shared-memory Programming
- Pthread
- Synchronization Problem & Tools
 - Pthread
 - ◆ Mutually exclusion Lock
 - ◆ Condition variable
 - POSIX Semaphore
 - JAVA Monitor
- Other issues

Critical Section & Mutual Exclusion

- **Critical Section** is a piece of code that can only be accessed by one process/thread at a time
- **Mutual exclusion** is the problem to insure only one process/thread can be in a critical section
- E.g.: The design of entry section & exit section provides mutual exclusion for the critical section



Locks

- Lock: the simplest mechanism for ensuring mutual exclusion of critical section

➤ Spinlock is one of the implementation:

```
while (lock == 1);           /* no operation in while loop */
lock = 1;                    /* enter critical section */
.
critical section
.
lock = 0;                    /* leave critical section */
```

- Locks are implemented in Pthreads by a special type of variables “mutex”
- **Mutex** is abbreviation of “mutual exclusion”

Pthread Lock/Mutex Routines

- To use mutex, it must be declared as of **type pthread_mutex_t** and initialized with **pthread_mutex_init()**
- A mutex is destroyed with **pthread_mutex_destroy()**
- A critical section can then be protected using **pthread_mutex_lock()** and **pthread_mutex_unlock()**
- Example:

```
#include "pthread.h"
pthread_mutex_t  mutex;
pthread_mutex_init (&mutex, NULL);
pthread_mutex_lock(&mutex);
Critical Section
pthread_mutex_unlock(&mutex);
pthread_mutex_destroy(&mutex);
```

specify default attribute for the mutex

// enter critical section

// leave critical section

Bounded-Buffer Problem


- A pool of n buffers, each capable of holding one item
- Producer:
 - grab an empty buffer
 - place an item into the buffer
 - waits if no empty buffer is available
- Consumer:
 - grab a buffer and retracts the item
 - place the buffer back to the free pool
 - waits if all buffers are empty

Bounded-Buffer Problem

- **Producer** process produces information that is consumed by a **Consumer** process

- Buffer as a circular array with size B

- next free: in
- first available: out
- empty: $in = out$
- full: $(in+1) \% B = out$



- The solution allows at most $(B-1)$ item in the buffer
 - Otherwise, cannot tell the buffer is full or empty


Shared-Memory Solution

```
/*producer*/
while (1) {
    while (((in + 1) % BUFFER_SIZE) == out)
        ; //wait if buffer is full
    buffer[in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
}
```


“*in*” only modified by producer

```
/*consumer*/
while (1) {
    while (in == out); //wait if buffer is empty
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
}
```

“*out*” only modified by consumer



```
/* global data structure */
#define BUFSIZE 10
item buffer[BUFSIZE];
int in = out = 0;
```



Using Mutex Lock

```
/*producer*/
```

```
while (1) {  
    nextItem = getItem( );  
    while (counter == BUFFER_SIZE) ;  
    buffer[in] = nextItem;  
    in = (in + 1) % BUFFER_SIZE;  
    mutex_lock(mutex);  
    counter++;  
    mutex_unlock(mutex);  
}
```

```
/*consumer*/
```

```
while (1) {  
    while (counter == 0) ;  
    item = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    mutex_lock(mutex);  
    counter--;  
    mutex_unlock(mutex);  
}
```

Condition Variables (CV)

- CV represent some **condition** that a thread can:
 - Wait on, until the condition occurs; or
 - Notify other waiting threads that the condition has occurred
- Three operations on condition variables:
 - **wait()** --- **Block** until another thread calls **signal()** or **broadcast()** on the CV
 - **signal()** --- Wake up **one thread** waiting on the CV
 - **broadcast()** --- Wake up **all threads** waiting on the CV
- In Pthread, CV **type** is a **pthread_cond_t**
 - Use **pthread_cond_init()** to initialize
 - **pthread_cond_wait (&theCV, &somelock)**
 - **pthread_cond_signal (&theCV)**
 - **pthread_cond_broadcast (&theCV)**

Using Condition Variable

■ Example:

- A threads is designed to **take action when x=0**
- Another thread is responsible for decrementing the counter

```
pthread_cond_t  cond;  
pthread_cond_init (cond, NULL);
```

```
pthread_mutex_t  mutex;  
pthread_mutex_init (mutex, NULL);
```

```
action() {  
    pthread_mutex_lock (&mutex)  
    if (x != 0)  
        pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

- All condition variable operation **MUST** be performed while a mutex is **locked!!!**

Semaphore

- A tool to generalize the synchronization problem

- **Deadlock** may occur if not use appropriately !

- More specifically...

- a record of **how many units** of a particular resource are available

- ◆ If #record = 1 ➔ **binary semaphore, mutex lock**

- ◆ If #record > 1 ➔ **counting semaphore**

- accessed only through 2 *atomic* ops: **wait** & **signal**

- **Spinlock** implementation:

- Semaphore is an **integer variable**

```
wait (S) {                               signal (S) {
    while (S <= 0) ;                      S++;
    S--;                                  }
}
```


Semaphore Example

- shared data:

semaphore S ; // initially $S = 1$

- Process P_i :

do {

 wait (S) ;

 critical section

 signal (S);

 remainder section

} while (1) ;

POSIX Semaphore

- Semaphore is part of **POSIX** standard BUT it is **not** belonged to Pthread
 - It can be used with or **without** thread
- POSIX Semaphore routines:
 - **sem_init**(sem_t *sem, int pshared, unsigned int value)
 - **sem_wait**(sem_t *sem)
 - **sem_post**(sem_t *sem)
 - **sem_getvalue**(sem_t *sem, int *valptr)
 - **sem_destroy**(sem_t *sem)

Initial value of the semaphore

Current value of the semaphore

- Example:

```
#include <semaphore.h>
sem_t sem;
sem_init(&sem);
sem_wait(&sem);
    // critical section
sem_post(&sem);
sem_destroy(&sem);
```

The Big Picture

- Getting **synchronization** right is hard!
- How to pick between locks, semaphores, convars, monitors???
- **Locks** are very **simple** for many cases
 - But may not be the most efficient solution
- **Condition variables** **allow threads to sleep** while holding a lock
 - Be aware whether they use Mesa or Hoare semantics
- **Semaphores** provide **general functionality**
 - But also make it really easy to mess up or cause deadlock

Reference

- Textbook:
 - Parallel Computing Chap8
- Pthread Tutorial
 - <https://computing.llnl.gov/tutorials/pthreads/>
- Synchronization Tools:
 - <http://www.eecs.harvard.edu/~mdw/course/cs61/mediawiki/images/7/7e/Lectures-semaphores.pdf>
- Pthread API:
 - <http://www.yolinux.com/TUTORIALS/LinuxTutorialPosixThreads.html>